

Package ‘textmatch’

November 30, 2019

Title Toolkit for Matching Textual Data and Evaluating Textual Similarity

Version 0.0.0.9000

Description What the package does (one paragraph).

Depends R (>= 3.5.2)

License What license is it under?

Encoding UTF-8

LazyData false

RoxygenNote 6.1.1.9000

Imports dplyr,
data.table,
quantda,
stats,
progress,
tm,
stm,
optmatch,
textir,
text2vec,
cem,
Rfast,
reshape2,
Matrix,
cobalt

Suggests knitr,
rmarkdown,
softmaxreg,
glmnet

VignetteBuilder knitr

R topics documented:

corp_contrast	2
FoxCNNcorpus	2
FoxCNNmeta	3
FoxCNNsurvey	3
get_bal_optmatch	4

get_CEM	5
get_matches	5
get_similarity_scores	6
get_word2vec_glove	6
pair_distances	7
quality_model	7
SE	8
select_matches	8
textPS	9
textPS_dist	9
transform_dfm	10
transform_stm	11
Index	12

corp_contrast	<i>Corpus contrasts</i>
---------------	-------------------------

Description

Corpus contrasts

Usage

corp_contrast(corp, Z)

Arguments

- corp a fitted [stm](#) object
- Z an indicator for treatment assignment

Value

Corpus contrast summaries

FoxCNNcorpus	<i>Corpus with 1,565 articles from CNN and 1,796 articles from Fox News.</i>
--------------	--

Description

Corpus with 1,565 articles from CNN and 1,796 articles from Fox News.

Usage

FoxCNNcorpus

Format

A [data.frame](#) with 3,361 observations of 5 articles

Details

Corpus of front-page news articles published online by CNN or Fox News from 12/20/2014 to 05/09/2015 containing 1,565 articles from CNN and 1,796 articles from Fox News. Data contains article identifiers corresponding to data in FoxCNNmeta as well as raw and cleaned text data.

References

Mozer et al. (2019) [Matching with Text Data: An Experimental Evaluation of Methods for Matching Documents and of Measuring Match Quality](#). *Political Analysis*, Forthcoming.

FoxCNNmeta	<i>Dataset containing metadata for FoxCNN corpus with 1,565 articles from CNN and 1,796 articles from Fox News.</i>
------------	---

Description

Dataset containing metadata for FoxCNN corpus with 1,565 articles from CNN and 1,796 articles from Fox News.

Usage

FoxCNNmeta

Format

A [data.frame](#) with 3,361 observations of 5 articles

Details

Metadata for the FoxCNN corpus including article names, original URLs, and dates of publication.

References

Mozer et al. (2019) [Matching with Text Data: An Experimental Evaluation of Methods for Matching Documents and of Measuring Match Quality](#). *Political Analysis*, Forthcoming.

FoxCNNsurvey	<i>Dataset containing distance measurements and average quality scores for sample of 505 pairs of matched Fox and CNN articles evaluated by human coders on Mechanical Turk. Raw quality scores are regressed on distance measurements to fit a predictive model for match quality as a function of the 117 metrics considered.</i>
--------------	---

Description

Dataset containing distance measurements and average quality scores for sample of 505 pairs of matched Fox and CNN articles evaluated by human coders on Mechanical Turk. Raw quality scores are regressed on distance measurements to fit a predictive model for match quality as a function of the 117 metrics considered.

Usage

```
FoxCNNsurvey
```

Format

A [data.frame](#) with 505 pairs and 108 features.

Details

Distance measurements and average quality scores for a sample of 505 matched pairs of documents evaluated by human coders.

References

Mozer et al. (2019) \"Matching with Text Data: An Experimental Evaluation of Methods for Matching Documents and of Measuring Match Quality\". *Political Analysis*, Forthcoming.

get_bal_optmatch	<i>Generates balance diagnostics for many possible matched datasets which can be used to make comparisons between matching methods.</i>
------------------	---

Description

Generates balance diagnostics for many possible matched datasets which can be used to make comparisons between matching methods.

Usage

```
get_bal_optmatch(matchobj, Z, covs)
```

Arguments

matchobj	an optmatch object or a matched dataset
Z	a vector of treatment indicators
covs	a matrix or data frame of covariates to assess balance on

Value

A [data.frame](#) MNIR sufficient reduction scores for a corpus

get_CEM	<i>Create a data frame of pairs of documents obtained through coarsened exact matching (CEM) within a specified number of bins and return indices for matched sets</i>
---------	--

Description

Create a data frame of pairs of documents obtained through coarsened exact matching (CEM) within a specified number of bins and return indices for matched sets

Usage

```
get_CEM(x, Z, rep.name, cuts, caliper_fun, verbose = FALSE, SR = NULL)
```

Arguments

x	a text representation
Z	a vector of treatment indicators
rep.name	a string or character with the name of the matching method
cuts	a function for how the variables will be binned. Defaults to "median"

Value

A [data.frame](#) of indices for matched pairs of documents

get_matches	<i>Similarity and distance computation between documents or features</i>
-------------	--

Description

These functions compute distance matrices from a text representation where each row is a document and each column is a feature to measure distance over based on treatment indicator Z

Usage

```
get_matches(dist, Z, dist.name, caliper_fun, within.calip = "quantile",
  tol = 0.001)
```

Arguments

Z	a vector of treatment indicators
dist.name	a string or character with the name of the matching method
caliper_fun	an optional function specifying the caliper to enforce when matching
x	a matrix of pairwise distances for all potential matches of treatment and control units. See pair_distances .

`get_similarity_scores` *This function calculates an input character vector's similarity matrix according to the measures contained in the predictive model.*

Description

This function calculates an input character vector's similarity matrix according to the measures contained in the predictive model.

Usage

```
get_similarity_scores(x)
```

Arguments

`x` A character vector where each element is a document

Value

A data frame of rows ($n * n - 1$) and columns 16; each column is one of the constituent similarity measures

`get_word2vec_glove` *This function calculates the Word2Vec embeddings*

Description

This function calculates the Word2Vec embeddings

Usage

```
get_word2vec_glove(dir.source, corpus)
```

Arguments

`dat` FoxCNN corpus to calculate Word2Vec scores for

Value

A list of data frames containing the Word2Vec projections of the corpus

pair_distances	<i>Similarity and distance computation between documents or features</i>
----------------	--

Description

These functions compute distance matrices from a text representation where each row is a document and each column is a feature to measure distance over based on treatment indicator Z

Usage

```
pair_distances(dat, Z, include = c("cosine", "euclidean", "mahalanobis"),
  form = "data.frame", verbose = FALSE)
```

Arguments

include	Which distances to calculate
form	Should the distances be returned as a list of matrices or condensed into a single data frame?
x	a matrix text representation with rows corresponding to each document in a corpus and columns that represent summary measures of the text (e.g., word counts, topic proportions, etc.). Acceptable forms include a valid quanteda dfm object, a tm Document-Term Matrix, a matrix of estimated topic proportions, or a vector of estimated propensity scores.

Value

A matrix showing pairwise distances for all potential matches of treatment and control units under various distance metrics

quality_model	<i>Fitted model for pairwise match quality as a function of 117 distance metrics calculated in Mozer et al. (2019). Trained on "FoxCNNsurvey" dataset.</i>
---------------	--

Description

Fitted model for pairwise match quality as a function of 117 distance metrics calculated in Mozer et al. (2019). Trained on "FoxCNNsurvey" dataset.

Usage

```
quality_model
```

Format

A **glmnet** model object.

Details

Fitted model for predicting the match quality score for a given pair of text documents as a function of 117 distance measurements.

References

Mozer et al. (2019) "Matching with Text Data: An Experimental Evaluation of Methods for Matching Documents and of Measuring Match Quality". *Political Analysis*, Forthcoming.

SE	<i>Utility functions for textmatch objects</i>
----	--

Description

Utility functions for textmatch objects

Usage

$SE(v, z)$

Arguments

v a vector of covariate or outcome data

Value

z a vector of treatment indicators

select_matches	<i>Given a model for match quality and a corpus of documents, calculate the estimated match quality of each potential pairing of treated and control documents and return a matched dataset containing all pairs of documents with estimated quality above a specified threshold</i>
----------------	--

Description

Given a model for match quality and a corpus of documents, calculate the estimated match quality of each potential pairing of treated and control documents and return a matched dataset containing all pairs of documents with estimated quality above a specified threshold

Usage

select_matches(corpus, Z, mod, threshold)

Arguments

corpus	description
Z	treatment indicator
threshold	for quality scores to return

Value

A [data.frame](#) of matched pairs of documents

textPS	<i>Calculate a univariate projection and pairwise distances for a corpus of text documents with binary treatment indicator Z.</i>
--------	---

Description

Calculate a univariate projection and pairwise distances for a corpus of text documents with binary treatment indicator Z.

Given a text representation x and treatment indicator Z, estimate a univariate projection of the text data that can serve as a sufficient statistic for the confounding information captured within the text.

Usage

```
textPS(x, Z, verbose = FALSE, g = 1)
```

```
textPS(x, Z, verbose = FALSE, g = 1)
```

Arguments

x	a TDM or DFM text representation
Z	a vector of treatment indicators
normalize	should the TDM features be normalized? Defaults to TRUE.
return.df	Should the distances be returned as a data frame? Default is TRUE.
x	a TDM or DFM text representation
Z	a vector of treatment indicators

Value

A [data.frame](#) MNIR sufficient reduction scores for a corpus

A [data.frame](#) MNIR sufficient reduction scores for a corpus

textPS_dist	<i>Given a text representation x and treatment indicator Z, construct a distance matrix for the pairwise distances between treatment and control documents based on the textual propensity score.</i>
-------------	---

Description

Given a text representation x and treatment indicator Z, construct a distance matrix for the pairwise distances between treatment and control documents based on the textual propensity score.

Usage

```
textPS_dist(x, Z, normalize = TRUE, return.df = TRUE,
  verbose = FALSE)
```

Arguments

x	a TDM or DFM text representation
Z	a vector of treatment indicators
normalize	should the TDM features be normalized? Defaults to TRUE.
return.df	Should the distances be returned as a data frame? Default is TRUE.

transform_dfm	<i>Applies bounds, weights, and/or coarsening schemes to a dfm or document frequency matrix to reduce the dimension of the data, reduce noise, or apply other design rules (e.g. - to exclude words that occur in too few or too many documents).</i>
---------------	---

Description

Applies bounds, weights, and/or coarsening schemes to a dfm or document frequency matrix to reduce the dimension of the data, reduce noise, or apply other design rules (e.g. - to exclude words that occur in too few or too many documents).

Usage

```
transform_dfm(x, bounds, tfidf = FALSE, verbose = TRUE)
```

Arguments

x	a matrix text representation with rows corresponding to each document in a corpus and columns that represent summary measures of the text (e.g., word counts, topic proportions, etc.). Acceptable forms include a valid quanteda dfm object, a tm Document-Term Matrix, or a matrix of estimated topic proportions.
bounds	a vector of lower and upper bounds to enforce. Defaults to excluding any terms that appear in only one document and any terms that appear in every document
tfidf	optional scheme to use for weighting the DTM. Defaults to FALSE.
verbose	indicator for verbosity

Value

A bounded DFM

transform_stm	<i>Refits a STM with a content-based covariate so that all document-level topic-proportions are estimated "as-treated". Also allows for calculation of the SR sufficient reduction and optional coarsening to reduce the dimension of the data, reduce noise, or apply other design rules (e.g. - to exclude words that occur in too few or too many documents).</i>
---------------	--

Description

Refits a STM with a content-based covariate so that all document-level topic-proportions are estimated "as-treated". Also allows for calculation of the SR sufficient reduction and optional coarsening to reduce the dimension of the data, reduce noise, or apply other design rules (e.g. - to exclude words that occur in too few or too many documents).

Usage

```
transform_stm(mod, out, Z, calc.SR = FALSE, coarsen = FALSE,  
             simplex = FALSE)
```

Arguments

mod	a fitted stm object
out	the original call to the STM
Z	an indicator for treatment assignment
calc.SR	an indicator for returning the sufficient reduction. Default is TRUE.
coarsen	an indicator for returning the coarsened STM

Value

A bounded DFM

Index

- *Topic **datasets**
 - quality_model, [7](#)
- *Topic **data**
 - FoxCNNcorpus, [2](#)
 - FoxCNNmeta, [3](#)
 - FoxCNNsurvey, [3](#)
- corp_contrast, [2](#)
- data.frame, [2–5](#), [9](#)
- FoxCNNcorpus, [2](#)
- FoxCNNmeta, [3](#)
- FoxCNNsurvey, [3](#)
- get_bal_optmatch, [4](#)
- get_CEM, [5](#)
- get_matches, [5](#)
- get_similarity_scores, [6](#)
- get_word2vec_glove, [6](#)
- glmnet, [7](#)
- optmatch, [4](#)
- pair_distances, [5](#), [7](#)
- quality_model, [7](#)
- SE, [8](#)
- select_matches, [8](#)
- stm, [2](#), [11](#)
- textPS, [9](#)
- textPS_dist, [9](#)
- transform_dfm, [10](#)
- transform_stm, [11](#)